

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:00:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221105304, and current time: 1487221165305

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221105304&finishedTimeEnd=1487221165305>

02-16-2017 00:00:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221105304&finishedTimeEnd=1487221165305>

02-16-2017 00:00:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 0

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:01:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221165306, and current time: 1487221225307

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221165306&finishedTimeEnd=1487221225307>

02-16-2017 00:01:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221165306&finishedTimeEnd=1487221225307>

02-16-2017 00:01:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 0

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : AnalysisProvider updating its Authenticate Token...

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:02:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221225308, and current time: 1487221285421

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221225308&finishedTimeEnd=1487221285421>

02-16-2017 00:02:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221225308&finishedTimeEnd=1487221285421>

02-16-2017 00:02:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 0

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221285422, and current time: 1487221345423

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221285422&finishedTimeEnd=1487221345423>

02-16-2017 00:03:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is

<http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221285422&finishedTimeEnd=1487221345423>

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 10

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143578

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143579

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143582

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143583

02-16-2017 00:03:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143581

02-16-2017 00:03:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143584

02-16-2017 00:03:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143590

02-16-2017 00:03:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143591

02-16-2017 00:03:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143589

02-16-2017 00:03:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143592

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221345424, and current time: 1487221405426

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is

<http://y/ypd269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221345424&finishedTimeEnd=1487221405426>

02-16-2017 00:04:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is

<http://y/ypd269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221345424&finishedTimeEnd=1487221405426>

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 15

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing SPARK application_1486843207585_143511

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143586

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143587

02-16-2017 00:04:25 ERROR com.linkedin.drelephant.ElephantRunner :

02-16-2017 00:04:25 ERROR com.linkedin.drelephant.ElephantRunner :
java.security.PrivilegedActionException: java.net.ConnectException: Connection refused

at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:356)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1689)
at com.linkedin.drelephant.security.HadoopSecurity.doAs(HadoopSecurity.java:99)
at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:99)
at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:48)
at com.linkedin.drelephant.analysis.AnalyticJob.getAnalysis(AnalyticJob.java:232)
at com.linkedin.drelephant.ElephantRunner\$ExecutorThread.run(ElephantRunner.java:151)
at java.util.concurrent.Executors\$RunnableAdapter.call(Executors.java:471)
at java.util.concurrent.FutureTask.run(FutureTask.java:262)
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)
at java.lang.Thread.run(Thread.java:745)

Caused by: java.net.ConnectException: Connection refused

at java.net.PlainSocketImpl.socketConnect(Native Method)
at java.net.AbstractPlainSocketImpl.doConnect(AbstractPlainSocketImpl.java:339)

at java.net.AbstractPlainSocketImpl.connectToAddress(AbstractPlainSocketImpl.java:198)
at java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:182)
at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
at java.net.Socket.connect(Socket.java:579)
at sun.net.NetworkClient.doConnect(NetworkClient.java:175)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:432)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:527)
at sun.net.www.http.HttpClient.<init>(HttpClient.java:211)
at sun.net.www.http.HttpClient.New(HttpClient.java:308)
at sun.net.www.http.HttpClient.New(HttpClient.java:326)
at
sun.net.www.protocol.http.HttpURLConnection.getNewHttpClient(HttpURLConnection.java:998)
at sun.net.www.protocol.http.HttpURLConnection.plainConnect(HttpURLConnection.java:934)
at sun.net.www.protocol.http.HttpURLConnection.connect(HttpURLConnection.java:852)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:686)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:638)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:711)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)
at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:415)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:1436)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:312)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getAuthParameters(WebHdfsFileSystem.java:524)

at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.toUrl(WebHdfsFileSystem.java:545)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractFsPathRunner.getUrl(WebHdfsFileSystem.java:801)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:415)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getHdfsFileStatus(WebHdfsFileSystem.java:948)

at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getFileStatus(WebHdfsFileSystem.java:963)

at org.apache.hadoop.fs.FileSystem.exists(FileSystem.java:1424)

at
org.apache.spark.deploy.history.SparkFSFetcher.org\$apache\$spark\$deploy\$history\$SparkFSFetcher\$\$isLegacyLogDirectory(SparkFSFetcher.scala:186)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:143)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:99)

... 13 more

02-16-2017 00:04:25 ERROR com.linkedin.drelephant.ElephantRunner : Add analytic job id [application_1486843207585_143511] into the retry list.

02-16-2017 00:04:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143610

02-16-2017 00:04:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143611

02-16-2017 00:04:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143608

02-16-2017 00:04:27 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143609

02-16-2017 00:04:28 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143615

02-16-2017 00:04:28 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143613

02-16-2017 00:04:29 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143616

02-16-2017 00:04:29 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143617

02-16-2017 00:04:30 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing SPARK application_1486843207585_143593

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner :

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner :
java.security.PrivilegedActionException: java.net.ConnectException: Connection refused

- at java.security.AccessController.doPrivileged(Native Method)
- at javax.security.auth.Subject.doAs(Subject.java:356)
- at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1689)
- at com.linkedin.drelephant.security.HadoopSecurity.doAs(HadoopSecurity.java:99)
- at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:99)
- at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:48)
- at com.linkedin.drelephant.analysis.AnalyticJob.getAnalysis(AnalyticJob.java:232)
- at com.linkedin.drelephant.ElephantRunner\$ExecutorThread.run(ElephantRunner.java:151)
- at java.util.concurrent.Executors\$RunnableAdapter.call(Executors.java:471)
- at java.util.concurrent.FutureTask.run(FutureTask.java:262)

at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)
at java.lang.Thread.run(Thread.java:745)

Caused by: java.net.ConnectException: Connection refused

at java.net.PlainSocketImpl.socketConnect(Native Method)
at java.net.AbstractPlainSocketImpl.doConnect(AbstractPlainSocketImpl.java:339)
at java.net.AbstractPlainSocketImpl.connectToAddress(AbstractPlainSocketImpl.java:198)
at java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:182)
at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
at java.net.Socket.connect(Socket.java:579)
at sun.net.NetworkClient.doConnect(NetworkClient.java:175)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:432)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:527)
at sun.net.www.http.HttpClient.<init>(HttpClient.java:211)
at sun.net.www.http.HttpClient.New(HttpClient.java:308)
at sun.net.www.http.HttpClient.New(HttpClient.java:326)
at
sun.net.www.protocol.http.HttpURLConnection.getNewHttpClient(HttpURLConnection.java:998)
at sun.net.www.protocol.http.HttpURLConnection.plainConnect(HttpURLConnection.java:934)
at sun.net.www.protocol.http.HttpURLConnection.connect(HttpURLConnection.java:852)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:686)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:638)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:711)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)


```
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner$1.run(WebHdfsFileSystem.java:588)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:415)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner.run(WebHdfsFileSystem.java:584)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:1436)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:312)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getAuthParameters(WebHdfsFileSystem.java:524)
    at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.toUrl(WebHdfsFileSystem.java:545)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractFsPathRunner.getUrl(WebHdfsFileSystem.java:801)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:709)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner.access$100(WebHdfsFileSystem.java:559)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner$1.run(WebHdfsFileSystem.java:588)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:415)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem$AbstractRunner.run(WebHdfsFileSystem.java:584)
    at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getHdfsFileStatus(WebHdfsFileSystem.java:948)
    at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getFileStatus(WebHdfsFileSystem.java:963)
    at org.apache.hadoop.fs.FileSystem.exists(FileSystem.java:1424)
```

at
org.apache.spark.deploy.history.SparkFSFetcher.org\$apache\$spark\$deploy\$history\$SparkFSFetcher\$\$isLegacyLogDirectory(SparkFSFetcher.scala:186)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:143)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:99)

... 13 more

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner : Add analytic job id [application_1486843207585_143593] into the retry list.

02-16-2017 00:04:30 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing SPARK application_1486843207585_143600

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner :

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner :
java.security.PrivilegedActionException: java.net.ConnectException: Connection refused

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:356)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1689)

at com.linkedin.drelephant.security.HadoopSecurity.doAs(HadoopSecurity.java:99)

at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:99)

at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:48)

at com.linkedin.drelephant.analysis.AnalyticJob.getAnalysis(AnalyticJob.java:232)

at com.linkedin.drelephant.ElephantRunner\$ExecutorThread.run(ElephantRunner.java:151)

at java.util.concurrent.Executors\$RunnableAdapter.call(Executors.java:471)

at java.util.concurrent.FutureTask.run(FutureTask.java:262)

at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)

at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)

at java.lang.Thread.run(Thread.java:745)

Caused by: java.net.ConnectException: Connection refused

at java.net.PlainSocketImpl.socketConnect(Native Method)

at java.net.AbstractPlainSocketImpl.doConnect(AbstractPlainSocketImpl.java:339)

at java.net.AbstractPlainSocketImpl.connectToAddress(AbstractPlainSocketImpl.java:198)
at java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:182)
at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
at java.net.Socket.connect(Socket.java:579)
at sun.net.NetworkClient.doConnect(NetworkClient.java:175)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:432)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:527)
at sun.net.www.http.HttpClient.<init>(HttpClient.java:211)
at sun.net.www.http.HttpClient.New(HttpClient.java:308)
at sun.net.www.http.HttpClient.New(HttpClient.java:326)
at
sun.net.www.protocol.http.HttpURLConnection.getNewHttpClient(HttpURLConnection.java:998)
at sun.net.www.protocol.http.HttpURLConnection.plainConnect(HttpURLConnection.java:934)
at sun.net.www.protocol.http.HttpURLConnection.connect(HttpURLConnection.java:852)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:686)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:638)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:711)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)
at java.security.AccessController.doPrivileged(Native Method)
at javax.security.auth.Subject.doAs(Subject.java:415)
at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:1436)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:312)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getAuthParameters(WebHdfsFileSystem.java:524)

at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.toUrl(WebHdfsFileSystem.java:545)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractFsPathRunner.getUrl(WebHdfsFileSystem.java:801)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:415)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getHdfsFileStatus(WebHdfsFileSystem.java:948)

at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getFileStatus(WebHdfsFileSystem.java:963)

at org.apache.hadoop.fs.FileSystem.exists(FileSystem.java:1424)

at
org.apache.spark.deploy.history.SparkFSFetcher.org\$apache\$spark\$deploy\$history\$SparkFSFetcher\$\$isLegacyLogDirectory(SparkFSFetcher.scala:186)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:143)

at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:99)

... 13 more

02-16-2017 00:04:30 ERROR com.linkedin.drelephant.ElephantRunner : Add analytic job id [application_1486843207585_143600] into the retry list.

02-16-2017 00:04:30 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143606

02-16-2017 00:04:31 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143607

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The list of RM IDs are rm1,rm2

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Checking RM URL: <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/info>

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : yldap269.kmdc.att.com:8088 is ACTIVE

02-16-2017 00:05:25 INFO com.linkedin.drelephant.ElephantRunner : Fetching analytic job list...

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : Fetching recent finished application runs between last time: 1487221405427, and current time: 1487221465428

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The succeeded apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=SUCCEEDED&finishedTimeBegin=1487221405427&finishedTimeEnd=1487221465428>

02-16-2017 00:05:25 INFO com.linkedin.drelephant.analysis.AnalyticJobGeneratorHadoop2 : The failed apps URL is <http://yldap269.kmdc.att.com:8088/ws/v1/cluster/apps?finalStatus=FAILED&finishedTimeBegin=1487221405427&finishedTimeEnd=1487221465428>

02-16-2017 00:05:25 INFO com.linkedin.drelephant.ElephantRunner : Job queue size is 18

02-16-2017 00:05:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143441

02-16-2017 00:05:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143623

02-16-2017 00:05:25 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143642

02-16-2017 00:05:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143643

02-16-2017 00:05:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143626

02-16-2017 00:05:26 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143627

02-16-2017 00:05:27 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143624

02-16-2017 00:05:28 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143625

02-16-2017 00:05:29 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143631

02-16-2017 00:05:29 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing MAPREDUCE application_1486843207585_143628

02-16-2017 00:05:36 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143629

02-16-2017 00:05:37 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143634

02-16-2017 00:05:38 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143635

02-16-2017 00:05:38 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 3 analyzing MAPREDUCE application_1486843207585_143636

02-16-2017 00:05:38 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 2 analyzing MAPREDUCE application_1486843207585_143637

02-16-2017 00:05:38 INFO com.linkedin.drelephant.ElephantRunner : Executor thread 1 analyzing SPARK application_1486843207585_143511

02-16-2017 00:05:38 ERROR com.linkedin.drelephant.ElephantRunner :

02-16-2017 00:05:38 ERROR com.linkedin.drelephant.ElephantRunner :
java.security.PrivilegedActionException: java.net.ConnectException: Connection refused

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:356)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1689)

at com.linkedin.drelephant.security.HadoopSecurity.doAs(HadoopSecurity.java:99)

at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:99)

at org.apache.spark.deploy.history.SparkFSFetcher.fetchData(SparkFSFetcher.scala:48)

at com.linkedin.drelephant.analysis.AnalyticJob.getAnalysis(AnalyticJob.java:232)

at com.linkedin.drelephant.ElephantRunner\$ExecutorThread.run(ElephantRunner.java:151)

at java.util.concurrent.Executors\$RunnableAdapter.call(Executors.java:471)
at java.util.concurrent.FutureTask.run(FutureTask.java:262)
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)
at java.lang.Thread.run(Thread.java:745)

Caused by: java.net.ConnectException: Connection refused

at java.net.PlainSocketImpl.socketConnect(Native Method)
at java.net.AbstractPlainSocketImpl.doConnect(AbstractPlainSocketImpl.java:339)
at java.net.AbstractPlainSocketImpl.connectToAddress(AbstractPlainSocketImpl.java:198)
at java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:182)
at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
at java.net.Socket.connect(Socket.java:579)
at sun.net.NetworkClient.doConnect(NetworkClient.java:175)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:432)
at sun.net.www.http.HttpClient.openServer(HttpClient.java:527)
at sun.net.www.http.HttpClient.<init>(HttpClient.java:211)
at sun.net.www.http.HttpClient.New(HttpClient.java:308)
at sun.net.www.http.HttpClient.New(HttpClient.java:326)
at
sun.net.www.protocol.http.HttpURLConnection.getNewHttpClient(HttpURLConnection.java:998)
at sun.net.www.protocol.http.HttpURLConnection.plainConnect(HttpURLConnection.java:934)
at sun.net.www.protocol.http.HttpURLConnection.connect(HttpURLConnection.java:852)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:686)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.connect(WebHdfsFileSystem.java:638)
at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:711)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:415)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:1436)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getDelegationToken(WebHdfsFileSystem.java:312)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getAuthParameters(WebHdfsFileSystem.java:524)

at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.toUrl(WebHdfsFileSystem.java:545)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractFsPathRunner.getUrl(WebHdfsFileSystem.java:801)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.runWithRetry(WebHdfsFileSystem.java:709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.access\$100(WebHdfsFileSystem.java:559)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner\$1.run(WebHdfsFileSystem.java:588)

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:415)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1709)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem\$AbstractRunner.run(WebHdfsFileSystem.java:584)

at
org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getHdfsFileStatus(WebHdfsFileSystem.java:948)
at org.apache.hadoop.hdfs.web.WebHdfsFileSystem.getFileStatus(WebHdfsFileSystem.java:963)
at org.apache.hadoop.fs.FileSystem.exists(FileSystem.java:1424)
at
org.apache.spark.deploy.history.SparkFSFetcher.org\$apache\$spark\$deploy\$history\$SparkFSFetcher\$\$is
LegacyLogDirectory(SparkFSFetcher.scala:186)
at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:143)
at org.apache.spark.deploy.history.SparkFSFetcher\$\$anon\$1.run(SparkFSFetcher.scala:99)
... 13 more